Brain Surgery for AI: How We Halved the Network Without Losing Accuracy

Big neural networks are smart — but often bloated. They carry around more neurons than they actually need to get the job done. This not only wastes computation but also makes them slower and harder to run on smaller devices.

In this paper, we introduce a method to **trim down neural networks by tracking which neurons actually "light up"** during prediction. By tracing the activation path for each class (we will go through MNIST dataset digits), we keep only the neurons that are truly doing the work — and drop the rest.

The trick? We don't just guess which neurons to cut. We use a **dynamic threshold-finding algorithm** that adjusts pruning levels automatically to keep accuracy above a level you choose. Want 95% accuracy? The system finds the tightest, leanest version of the network that still hits that target.

What's exciting is that **you can cut out up to half the neurons** — or more — and still get almost the same results. It's like shrinking your brain but keeping your math skills intact.

This approach opens doors for smarter, faster, and more efficient AI — especially on devices where size and speed matter. And the best part? It's simple, effective, and doesn't need fancy retraining.

Introduction / Problem Statement

In the race toward Artificial General Intelligence (AGI), neural networks have exploded in size — from billions to trillions of parameters. But this race comes with a cost: **inflated compute, energy consumption, and privacy concerns**.

While scaling up has enabled stunning breakthroughs, not every use case needs a trillion-parameter model. In fact, many **sensitive and high-stakes applications** — think medical diagnostics, legal data processing, or confidential business analytics — require a different kind of intelligence: **lean, efficient, and private**.

In these scenarios, sending data to massive external models isn't acceptable. You need models that are:

- **Small enough** to run locally or securely

- **Accurate enough** to still deliver meaningful results
- And ideally, **transparent enough** to be trusted and understood

This paper introduces a practical method to get there: by tracing neuron activations and pruning those that stay silent, we can **cut down model size dramatically** — with **no retraining** and **no meaningful drop in accuracy**.

Proposed Solution

Most neurons in a trained neural network stay inactive during inference — especially for specific classes or tasks. Our method takes advantage of this by **analyzing which neurons are used** and **removing the rest**.

We trace neuron activations on a test set, calculate how often each neuron fires, and **prune those that stay mostly inactive** (below a threshold). The result is a smaller, faster subnetwork that still achieves the desired accuracy — without retraining or changing the architecture.

This approach:

- Reduces computational load
- Maintains or even improves accuracy
- Offers better control and privacy
- Requires no additional training

Conclusion & Next Steps

This paper introduced a simple yet powerful idea: large neural networks often use only a fraction of their capacity for specific tasks — and we can safely remove the unused parts without hurting performance.

The results speak for themselves:

- **Up to 60–80% of neurons pruned**
- **No loss in accuracy**
- **10 specialized subnetworks**, each tailored for a single digit class

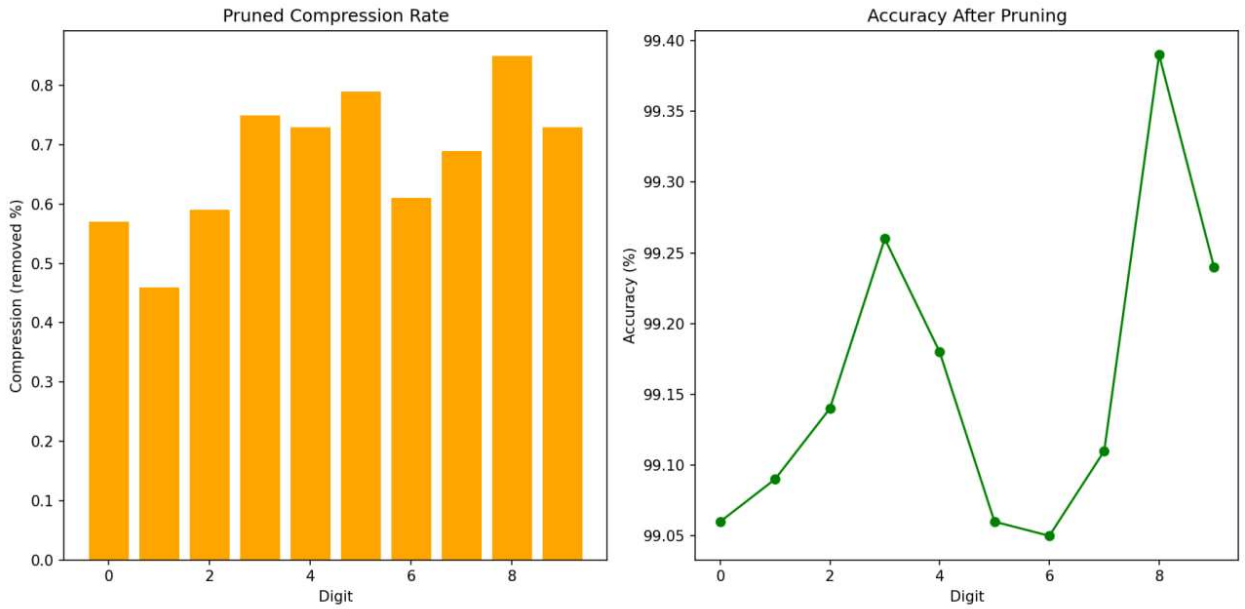

Check the GitHub Repository

If you're working on:

- Network compression
- On-device inference
- Privacy-focused AI
- Efficient models for narrow domains

...or if this sparks a use case in your work — I'd love to chat.

Feel free to reach out! - ubaydullohpulat@gmail.com or LinkedIn